

WEEK - 3

- How to work with data in R?
 - stored in R [data frame]
 - Read in data into R
- Simulate samples from a given distribution.

Recall - Week 1 and 2


- How to work with R and Rstudio?
- Data Visualisation :-
 - ggplot
 - plot ..

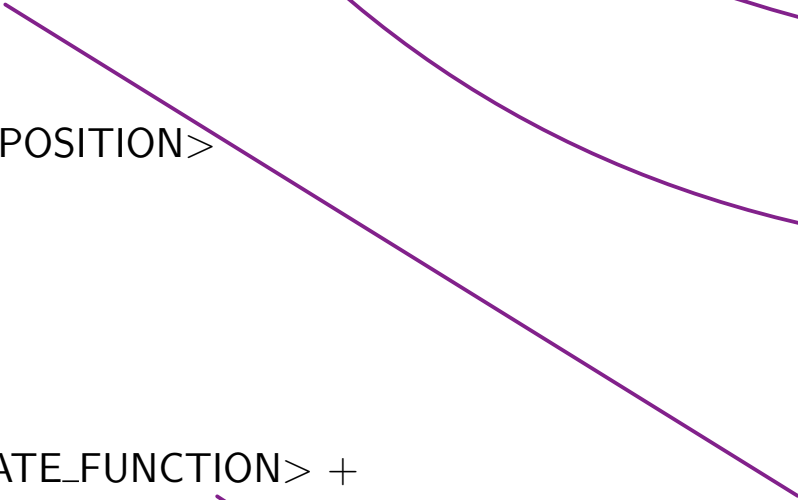
Recall :-

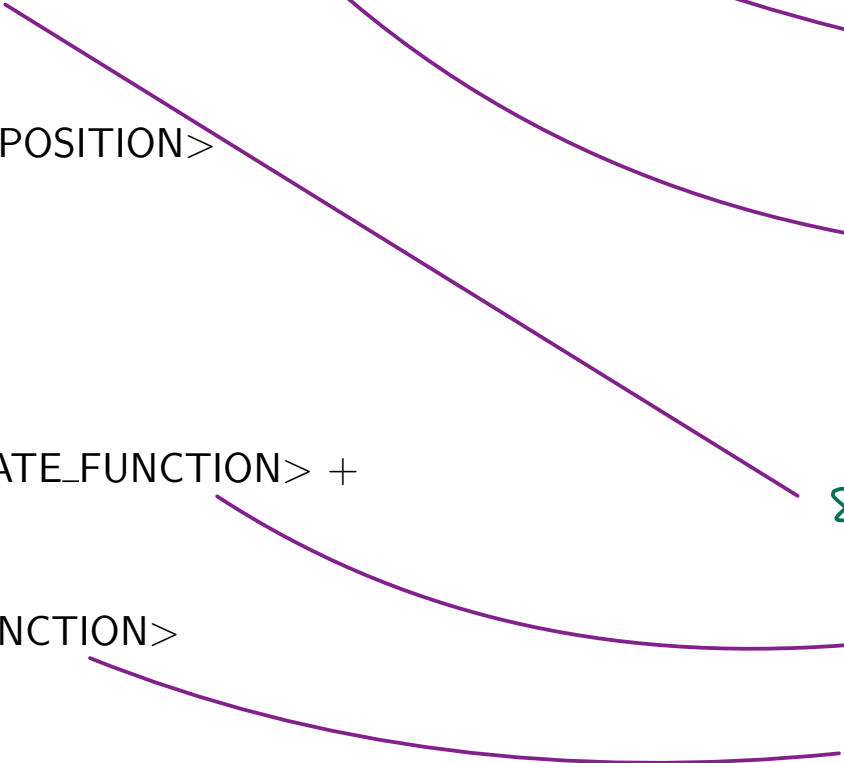
Layered Grammar of graphics

ggplot(data = <DATA>)

+ <GEOM_FUNCTION>(


mapping = aes(<MAPPINGS>),


stat = <STAT>,


position = <POSITION>


) +

<COORDINATE_FUNCTION> +

<FACET_FUNCTION>

- Each template takes
7 statements / Parameters

Specify data frame

geom_point, ... [Type of
geometry]
in graph]

map variables to an
aesthetic

statistical transformation

flip, polar
faceting

Data Types in R

```
> Course = "B.Sc."
```

```
> Number = 40
```

```
> Smart = TRUE
```

```
> mode(Course)
[1] "character"
```

```
> mode(Number)
[1] "numeric"
```

```
> mode(Smart)
[1] "logical"
```

R has many data types
- Focus on three

- character data = Surrounded by double quotes
- logical data = TRUE or FALSE
- numeric data
- Another important characteristic is class. For this we can use the class function.

Creating Vectors in R

```
> x = 3:7  
> x  
[1] 3 4 5 6 7
```

Vectors - variables

```
> s = seq(1,10, by=1)  
> s  
[1] 1 2 3 4 5 6 7 8 9 10
```

can be created by
seq -

```
> s10 = seq(1,10,by=0.5)  
> s10  
[1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5 6.0  
[12] 6.5 7.0 7.5 8.0 8.5 9.0 9.5 10.0
```

Position of 6.5 in the
vector

```
> rep(6,7)  
[1] 6 6 6 6 6 6 6  
> rep(x,3)  
[1] 3 4 5 6 7 3 4 5 6 7 3 4 5 6 7
```

Repeats the first argument
6, as many times as the
2nd argument 7.

first argument can be
a vector

Vectors in R

```
> KA_D = c(215,620,558, 1109,8813,350, 780, 420,  
+          144,478,816,242,1051,249,1238, 315,  
+          807,185,1993,515,1997,2886,371,156,  
+          589,1746,838,964,296,128)
```

```
> KA_D[c(1,3,5)]
```

```
[1] 215 558 8813
```

```
> KA_D[-c(1:20)]
```

```
[1] 1997 2886 371 156 589 1746 838 964 296 128
```

```
> KA_Dp = (KA_D)/sum(KA_D)*100
```

```
> KA_Dp
```

```
[1] 0.6964916 2.0084875 1.8076387 3.5926010 28.5496777
```

```
[6] 1.1338236 2.5268068 1.3605883 0.4664874 1.5484791
```

```
[11] 2.6434287 0.7839580 3.4047102 0.8066345 4.0104960
```

```
[16] 1.0204412 2.6142732 0.5993067 6.4563154 1.6683404
```

```
[21] 6.4692734 9.3491853 1.2018530 0.5053614 1.9080631
```

```
[26] 5.6561599 2.7146976 3.1228741 0.9588908 0.4146555
```

Karnataka Covid 19

Bulletin

• Entered it physically
into R - as a vector

• select

• Delete

• Apply a certain
function to the vector

Vectors in R

```
> KA_D[KA_D < 1000]
```

```
[1] 215 620 558 350 780 420 144 478 816 242 249 315 807 185  
[15] 515 371 156 589 838 964 296 128
```

```
> sum(KA_D > 2000)
```

```
[1] 2
```

```
> max(KA_D)
```

```
[1] 8813
```

```
> which(KA_D == max(KA_D))
```

```
[1] 5
```

logical operators can
be used to
perform selection
and identification

function max

Missing Values in R

```
> x = c(1,45,6,7,NA,99,0)
```

```
> x
```

```
[1] 1 45 6 7 NA 99 0
```

```
> x==NA
```

```
[1] NA NA NA NA NA NA NA
```

```
> mean(x)
```

```
[1] NA
```

```
> mean(x,na.rm=TRUE)
```

```
[1] 26.33333
```

```
> is.na(x)
```

```
[1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE
```

```
> mean(x[!is.na(x)])
```

```
[1] 26.33333
```

Missing values

- remove them before
performing / applying
functions on them

- logical operators can't
be used to identify
them

- `is.na` or `na.rm`
commands to identify.

Modes in R

```
> x = c("Siva", "looser", "3", "5")
```

```
> mode(x)  
[1] "character"
```

```
> x[3] + x[4]
```

```
> as.numeric(x[3]) + as.numeric(x[4])  
[1] 8
```

Vector:- each element is
in the same mode

Computation would lead to
an error.

R has to be told to
treat it as a
numeric before one
does the computation.

Matrices in R

```
> A = matrix(seq(3,5, by=1/10), 7,3)
```

```
> A
```

	[,1]	[,2]	[,3]
[1,]	3.0	3.7	4.4
[2,]	3.1	3.8	4.5
[3,]	3.2	3.9	4.6
[4,]	3.3	4.0	4.7
[5,]	3.4	4.1	4.8
[6,]	3.5	4.2	4.9
[7,]	3.6	4.3	5.0

All elements in the matrix have the same mode

vector used to create matrix.

Specify rows & columns

```
> B = matrix(seq(3,5, by=1/10), ncol=3)
```

fill the entries row wise.

```
> C = matrix(seq(3,5, by=1/10), ncol=3, byrow=TRUE)
```

```
> A[4,1]
```

```
[1] 3.3
```

Data Frames in R

```
> xd = c("Siva", "looser", 3, 5)
```

```
> xd  
[1] "Siva"  "looser" "3"      "5"
```

In R will convert
all elements of the vector
to the same mode

Data Frames in R

District names

```
> KA_District=c("Bagalakote","Ballari","Belagavi",  
+ "Bengaluru Rural","Bengaluru Urban", "Bidar","Chamarajanagara",  
+ "Chikkaballapura","Chikkamagaluru","Chitradurga", "Dakshina Kannada",  
+ "Davanagere","Dharwada","Gadag","Hassana","Haveri","Kalaburagi",  
+ "Kodagu","Kolara","Koppala","Mandya","Mysuru","Raichuru",  
+ "Ramanagara","Shivamogga","Tumakuru","Udupi","Uttara Kannada",  
+ "Vijayapura","Yadagiri")
```

Discharge Data

```
> KA_Discharge = data.frame(KA_District, KA_D)  
> class(KA_Discharge)  
[1] "data.frame"  
> mode(KA_Discharge)  
[1] "list"
```

sapply - simplifies
"loop function"

```
> sapply(KA_Discharge,mode)
```

```
KA_District      KA_D  
"character"      "numeric"
```

Data frame

- rectangular array

- elements can be in different mode

{
 . 1st column
 - character
 . 2nd column
 - numeric

Data Frames as Matrix in R

```
> names(KA_Discharge)=c("District", "Recovered")
```

change the names
of the variables in
the dataframe

```
> KA_Discharge$Recovered-
```

```
[1] 215 620 558 1109 8813 350 780 420 144 478 816  
[12] 242 1051 249 1238 315 807 185 1993 515 1997 2886  
[23] 371 156 589 1746 838 964 296 128
```

```
> KA_Discharge[3,2]
```

```
[1] 558
```

```
> KA_Discharge[3,]
```

```
District Recovered  
3 Belagavi 558
```

```
> KA_Discharge[, "Recovered"]
```

```
[1] 215 620 558 1109 8813 350 780 420 144 478 816  
[12] 242 1051 249 1238 315 807 185 1993 515 1997 2886  
[23] 371 156 589 1746 838 964 296 128
```

various ways
of selecting
objects in a
data frame

Data Frames as Matrix in R

```
> Deaths= c(346, 1712, 975, 903, 16593, 407,515, 446, 400,221,  
+ 1750, 611,1333,328,1291,652,856, 343, 647, 530, 673,  
+ 2494, 346, 338, 1105, 1172, 509, 793, 500, 206  
+ )
```

Death counts from
the bulletin

```
> KA_Discharge$Deaths = Deaths
```

- Created a new
variable
- name = Deaths

```
> head(KA_Discharge)
```

- assigned values
from vector created
earlier.

	District	Recovered	Deaths
1	Bagalakote	215	346
2	Ballari	620	1712
3	Belagavi	558	975
4	Bengaluru Rural	1109	903
5	Bengaluru Urban	8813	16593
6	Bidar	350	407

head() - displays
of the dataframe

Reading Data Frames into R

```
> kabbuldf=read.csv(file="KAbulletin.csv", header = TRUE)
```

```
> names(kabbuldf)
```

[1] "District"	"Today.s.Positives"
[3] "Total.Positives"	"Today.s.Discharges"
[5] "Total.Discharges"	"Total.Active.Cases"
[7] "Today.s.Deaths"	"Deaths"

- R/R-studio : please ensure you set the correct working directory

└ - where csv file is located.

Used to read csv files.

other read commands

- read.table

- read.xls ...

(require package)

1st row of csv file - provides names of variables

Reading Data Frames into R

```
> head(kabulldf)
```

	District	Today.s.Positives	Total.Positives
1	Bagalakote	325	39150
2	Ballari	502	111730
3	Belagavi	900	91717
4	Bengaluru Rural	517	77877
5	Bengaluru Urban	10692	1720890
6	Bidar	84	28865

	Today.s.Discharges	Total.Discharges	Total.Active.Cases
1	215	36541	2263
2	620	104268	5750
3	558	83958	6784
4	1109	72741	4233
5	8813	1570258	134038
6	350	27619	835

	Today.s.Deaths	Deaths
1	0	346
2	5	1712
3	6	975
4	0	903
5	12	16593
6	0	407

head()

- 8 variables

- Ex:

• mode (kabulldf)

• apply (kabulldf, mode)

Selecting from Data Frames in R

```
> kaballdf[which.max(kaballdf$"Today.s.Positives"),]
```

	District	Today.s.Positives	Total.Positives
32	Total	24172	3809467
	Today.s.Discharges	Total.Discharges	Total.Active.Cases
32	30869	3526108	244331
	Today.s.Deaths	Deaths	
32	56	38998	

Selects the district that has the maximum "Today.s.Positives"
☹️ Total is largest

```
> hpkaballdf = subset(kaballdf, 1st Today.s.Positives > 1000)
```

```
> head(hpkaballdf,2)
```

	District	Today.s.Positives	Total.Positives
5	Bengaluru Urban	10692	1720890
13	Dharwada	1044	80274
	Today.s.Discharges	Total.Discharges	Total.Active.Cases
5	8813	1570258	134038
13	1051	72514	6425
	Today.s.Deaths	Deaths	
5	12	16593	
13	1	1333	

Queries

- find objects that have certain properties

Two arguments

- 1st = data frame
- 2nd = condition that you want to use to create subset

Selecting from Data Frames in R

```
> IRDkabulldf = subset(kabulldf,  
+   select=c("Total.Positives", "Total.Discharges", "Deaths"))
```

1st

Data frame

Specifies the
columns to be
subset-ed.

2nd

```
> head(IRDkabulldf)
```

	Total.Positives	Total.Discharges	Deaths
1	39150	36541	346
2	111730	104268	1712
3	91717	83958	975
4	77877	72741	903
5	1720890	1570258	16593
6	28865	27619	407

Data frame

Created by subset
Command retains
the same
names for the
variables as the
original dataframe.

Ordering Data Frames in R

```
> okabulldf = kabulldf[order(kabulldf$Today.s.Positives),]
```

```
> head(okabulldf, 4)
```

	District	Today.s.Positives	Total.Positives
31	Others*	0	36
6	Bidar	84	28865
24	Ramanagara	84	29064
29	Vijayapura	121	39690

	Today.s.Discharges	Total.Discharges	Total.Active.Cases
31	0	33	0
6	350	27619	835
24	156	27559	1167
29	296	38080	1110

	Today.s.Deaths	Deaths
31	0	3
6	0	407
24	0	338
29	0	500

Re order the
rows of the dataframe
- corresponding to the
order of one
variable

① - variable =
Today.s.Positives

check: options in
order command

Generating Random data in R

$X \sim \text{Uniform}(\{1, 2, \dots, n\})$:

Let $n \geq 1$ be an integer. If X is a random variable such that

$$P(X = k) = \frac{1}{n} \text{ for all } 1 \leq k \leq n$$

```
> sample(1:6, 10, replace=T)
```

```
[1] 6 3 3 1 1 2 1 2 4 5
```

Rolling a dice 10 times
fair dice

```
> sample(c(0,1), 10, replace = TRUE, prob = c(0.3, 0.7))
```

Tossing a biased(0.7)
Coin 10 times

Goal :-

- Generate samples from a given distribution.

- n=6 - Experiment is rolling a fair dice

vector to sample from
{1, 2, 3, 4, 5, 6}

of samples to generate
10

Sample with replacement

{ - 0 with probability 0.3
1 with probability 0.7

Generating Random data in R

$X \sim \mathbf{Binomial}(n, p)$: Let $0 \leq p \leq 1$ and let $n \geq 1$ be an integer. If X is a random variable taking values in $\{0, 1, \dots, n\}$ having a probability mass function

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

for all $0 \leq k \leq n$.

```
> rbinom(10, 6, 0.5)
```

```
[1] 3 4 5 5 2 3 4 2 3 2
```

```
> rbinom(m, size, prob)
```

```
> rbinom(10, 30, 0.3)
```

```
[1] 13 11 8 8 8 7 11 9 6 11
```

$\left\{ \begin{array}{l} - n \text{ Bernoulli trials} \\ - (\text{Tossing a coin } n \text{ times}) \\ - \# \text{ of successes} \\ - (\# \text{ of Heads or 1's}) \end{array} \right.$

Binomial Experiment

of samples required

of Bernoulli trials (n)

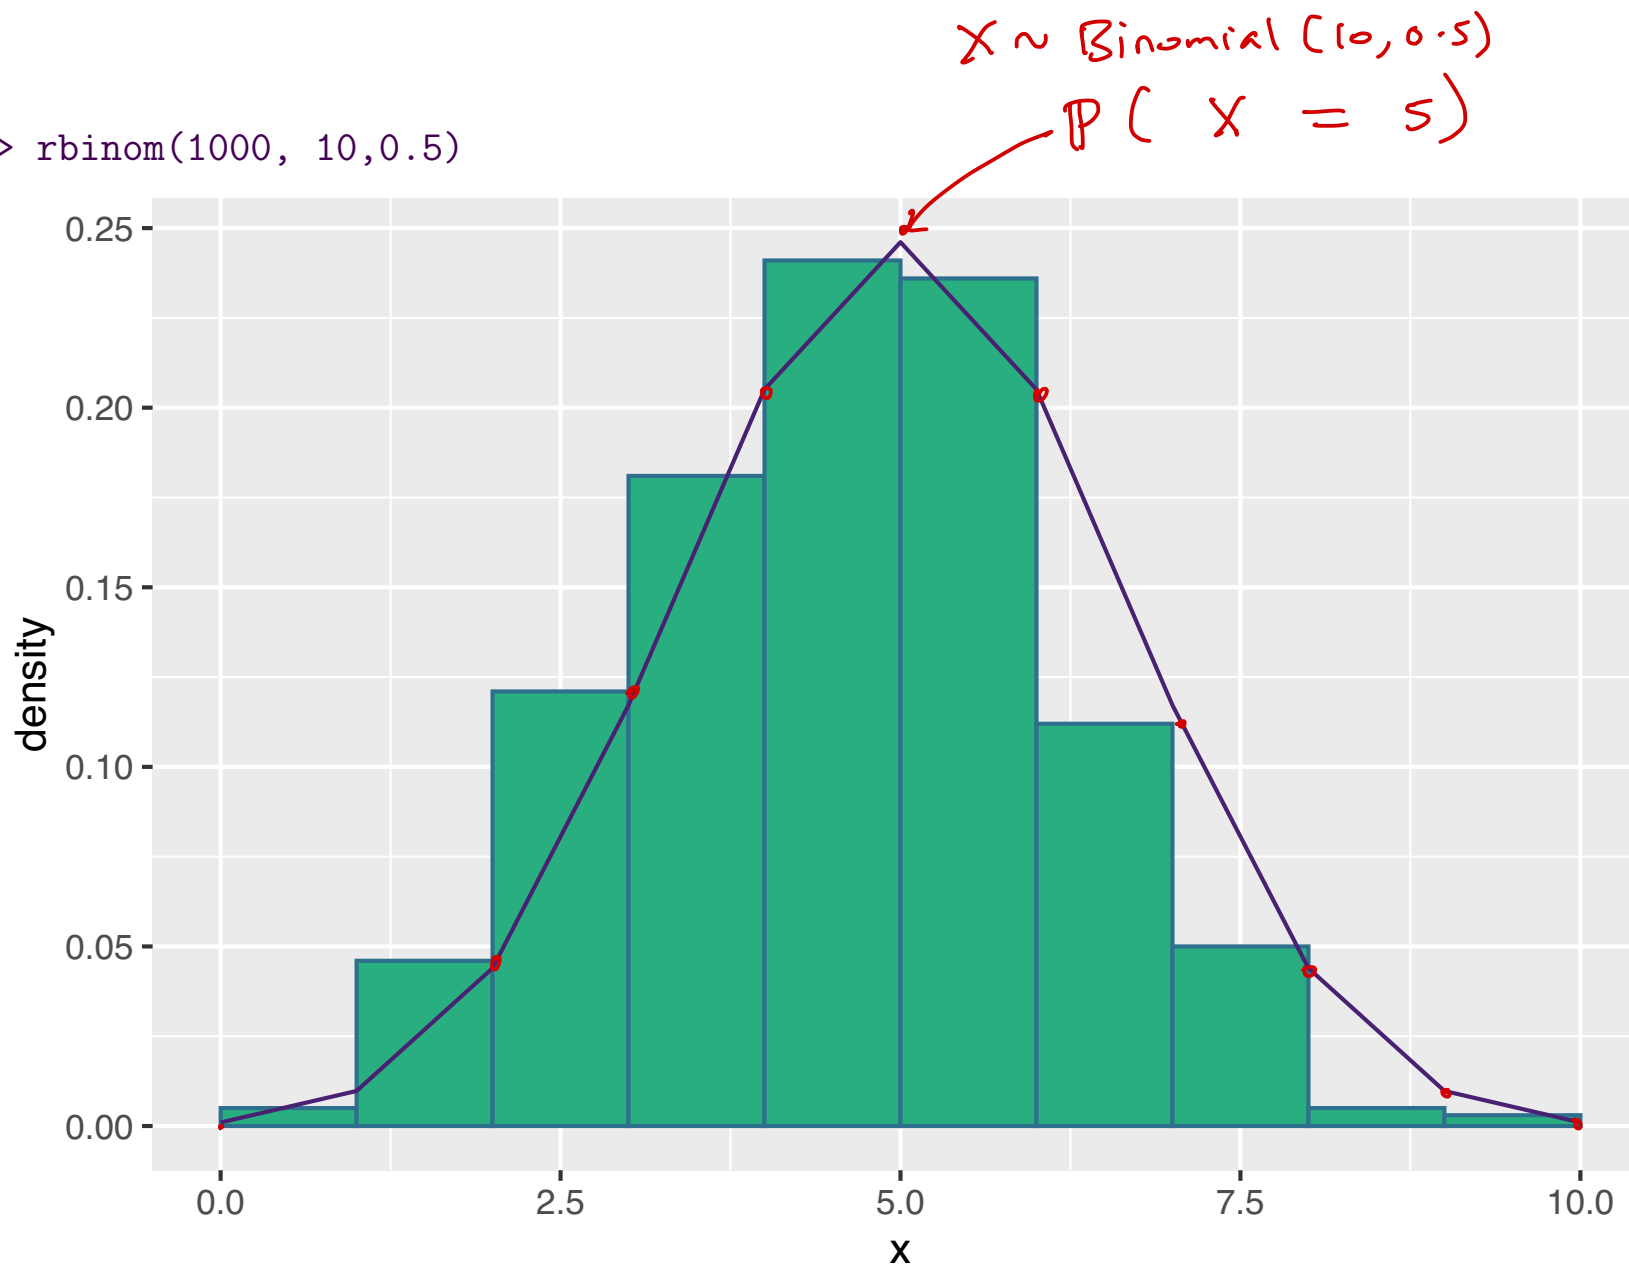
Prob of Success (p)

Ex:-
check

dbinom()
rbinom()
pbinom()

Generating Random data in R

```
> rbinom(1000, 10, 0.5)
```



Work sheet- Exercise

- Generate 1000 samples of Binomial(10, 0.5)
- Histogram of generated data (proportions)
- line plot of the true Binomial probabilities

Generating Random data in R

$X \sim \text{Normal}(\mu, \sigma^2)$: Let $\mu \in \mathbb{R}$ and let $\sigma > 0$. Then X is said to be normally distributed with parameters μ and σ^2 if it has the density

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad (1)$$

for all $x \in \mathbb{R}$.

```
> rnorm(1, 10, 5)
```

```
[1] 5.865734
```

```
> rnorm(n, mean, sd)
```

of samples
standard deviation
mean of normal random variable

```
> rnorm(10, 3, 5)
```


```
[1] 7.529100 7.972982 -7.149743 6.655984 2.105155
```

```
[6] 1.114047 9.126808 7.754853 7.459944 4.858321
```

Normal distribution

$\sigma=1, \mu=0$

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}}$$



Ex:-

check


`dnorm()`

`pnorm()`

`qnorm()`

Generating Random data in R

```
> rnorm(1000, 0,1)
```



Work sheet-

Exercise

- Generate 1000 samples of Normal (0,1)

Histogram of generated data
(proportions)

- line plot of the true normal density

Generating Random data in R

$X \sim \mathbf{Exp}(\lambda)$: Suppose $\lambda > 0$. If X is a random variable with its probability density function given by

$$f(x) = \begin{cases} \lambda e^{-\lambda x} & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

it is said to be distributed exponentially with parameter λ .

```
> rexp(10, 1/2500)
```

```
[1] 10091.04499 3826.35655 537.16355 1443.33491  
[5] 5863.52814 2368.87468 2256.52472 4008.94340  
[9] 1390.32837 83.46584
```

```
> rexp(n, rate)
```

of samples


rate

```
> rexp(10, 3)
```

```
[1] 0.35294149 0.19999367 0.44155155 0.35013407 0.19273441  
[6] 0.07736766 0.39486037 0.08252493 0.03475531 0.10853122
```

Exponential distribution

rate $\equiv \lambda$




Ex :

check `dexp()`
`pexp()`
`qexp()`

Generating Random data in R

```
> rexp(1000, 1)
```



Work sheet-

Exercise

- Generate 1000 samples of $\text{Exponential}(1)$

Histogram of generated data
(proportions)

- line plot of the true exponential density